



Sandia
National
Laboratories

Exceptional service in the national interest

Building Automated Proofs of Refinement Between State Machines and C

Samuel D. Pollard, Sandia National Labs

Frama-C Days

Maison de la Radio et de la Musique, Paris

14 June 2024, 11:30

SAND2024-07142C

Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia LLC, a wholly owned subsidiary of Honeywell International Inc. for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.



Overview

- Sandia National Labs is a US government research & development center
- Sandia develops software for high-consequence embedded control systems



Livermore, California site



Overview

- The systems are relatively simple
- The cost for error is very high
- Requirements relatively complex
- A good use case for formal methods



[Emergency Services Sector](#)



[Energy Sector](#)



[Financial Services Sector](#)



[Critical Manufacturing Sector](#)



[Dams Sector](#)



[Defense Industrial Base Sector](#)



[Information Technology](#)



[Nuclear Reactors, Materials,](#)



[Transportation Systems Sector](#)



[Chemical Sector](#)



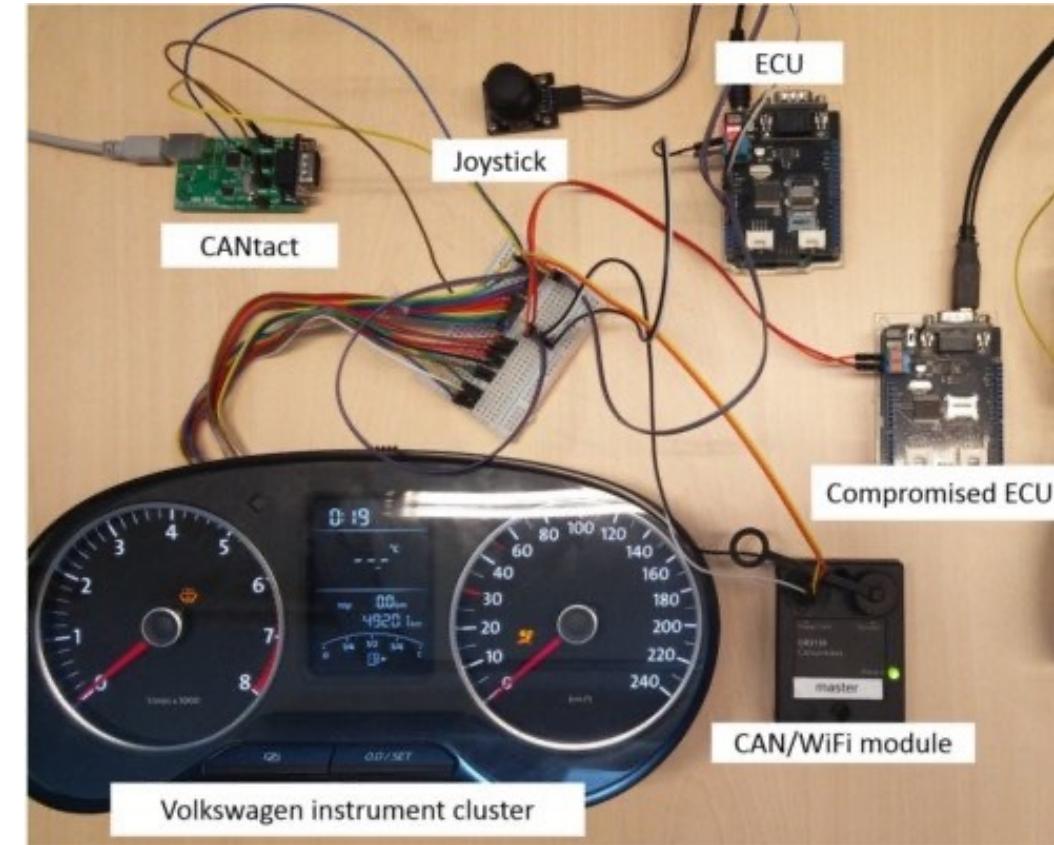
[Commercial Facilities Sector](#)



[Communications Sector](#)

Design Features of High Consequence Systems (HCS)

- Asynchronous interacting components
 - e.g., across a bus
- Requirements documents in English and informal diagrams
- Software implemented in C



From these, we require proofs of *system-level* properties

Introducing Q Framework

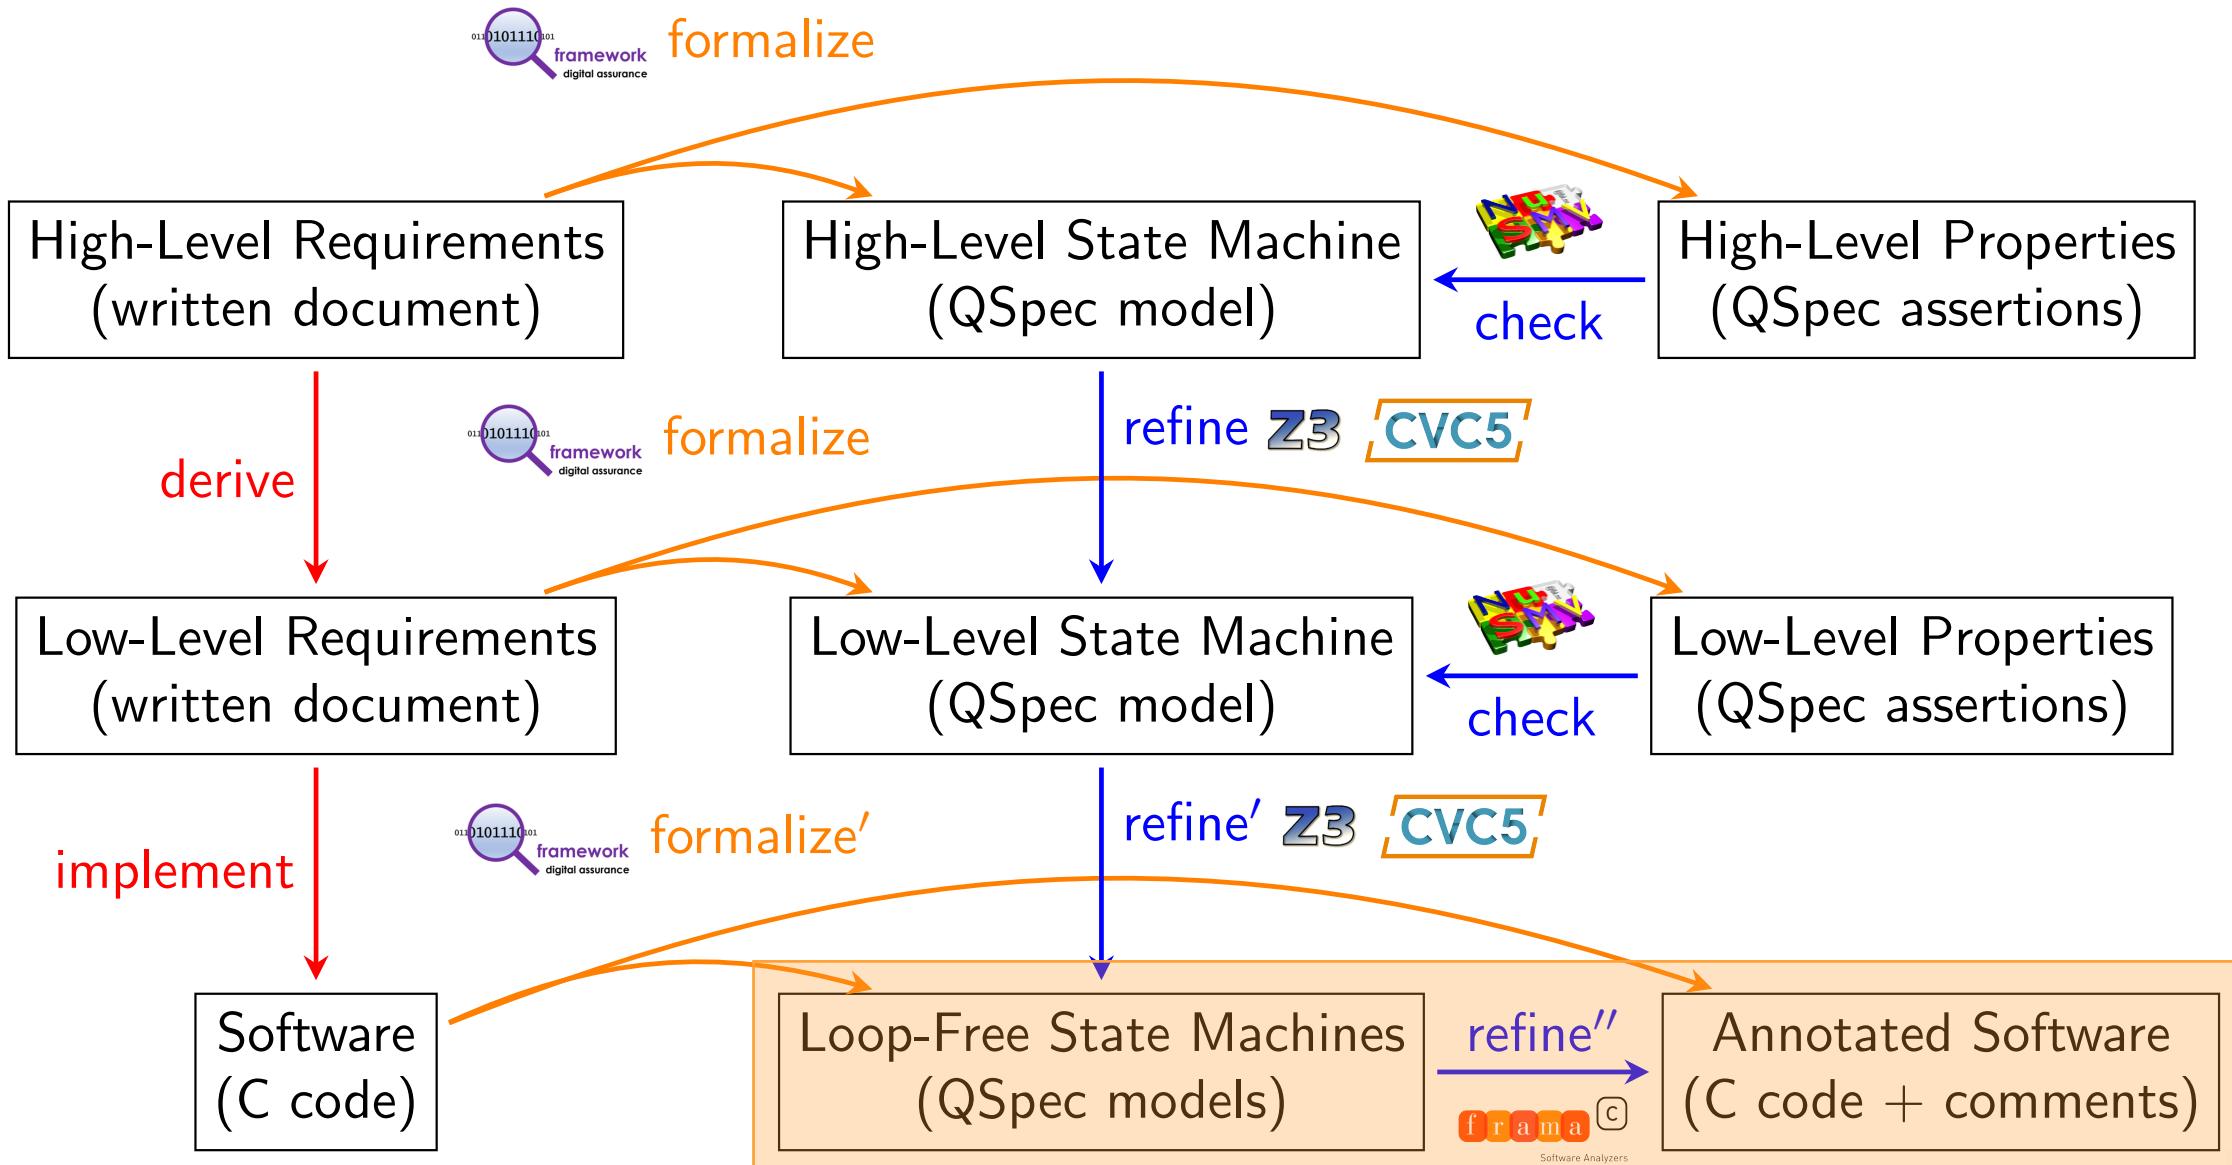
- Began in 2017
- Verify systems developed using model-based system design (MBSD)
- Leverage solvers for automation
 - NuSMV for LTL/CTL
 - Frama-C
- Currently has ~6 developers
- Part of a broader research group
 - hardware and software understanding
 - modeling, simulation, formal methods
- v1 in OCaml, v2 Haskell



Software Analyzers

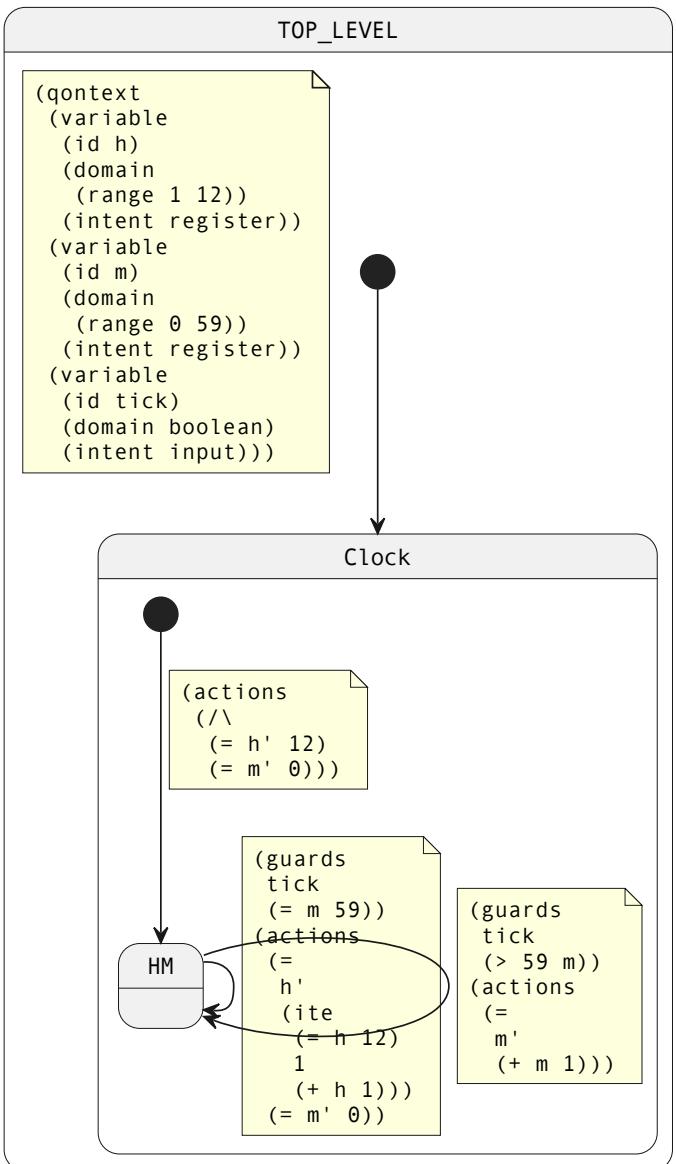


System-Level Refinement

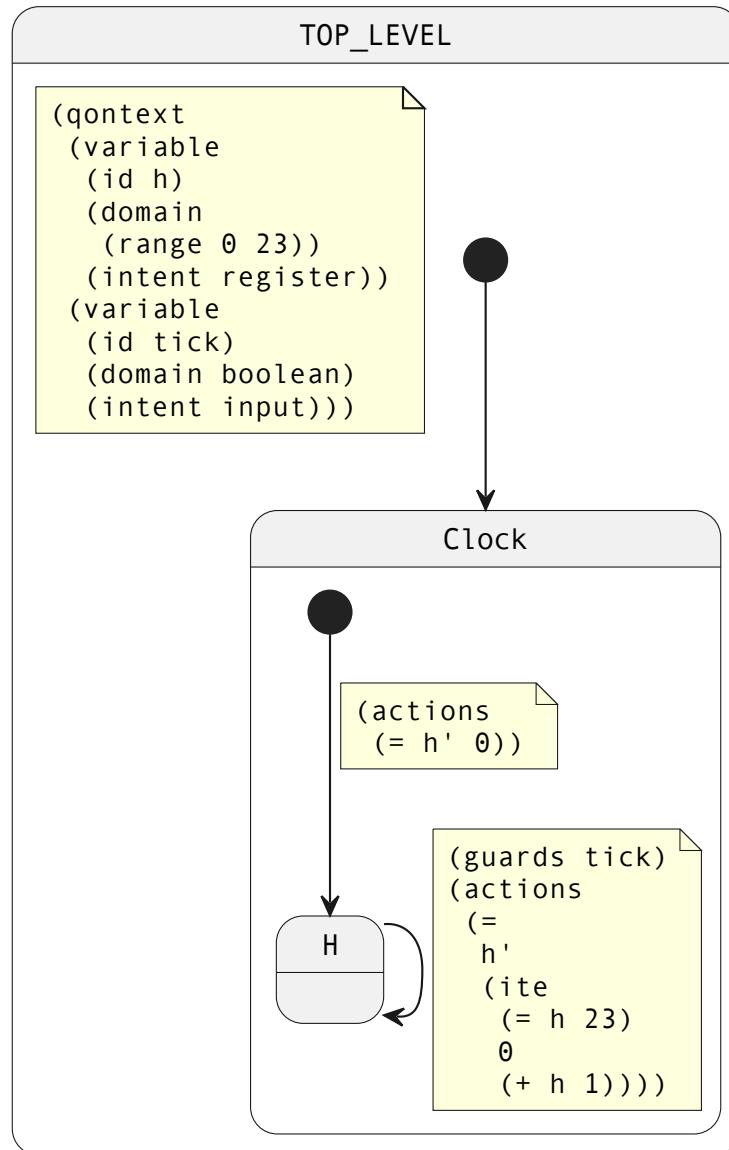




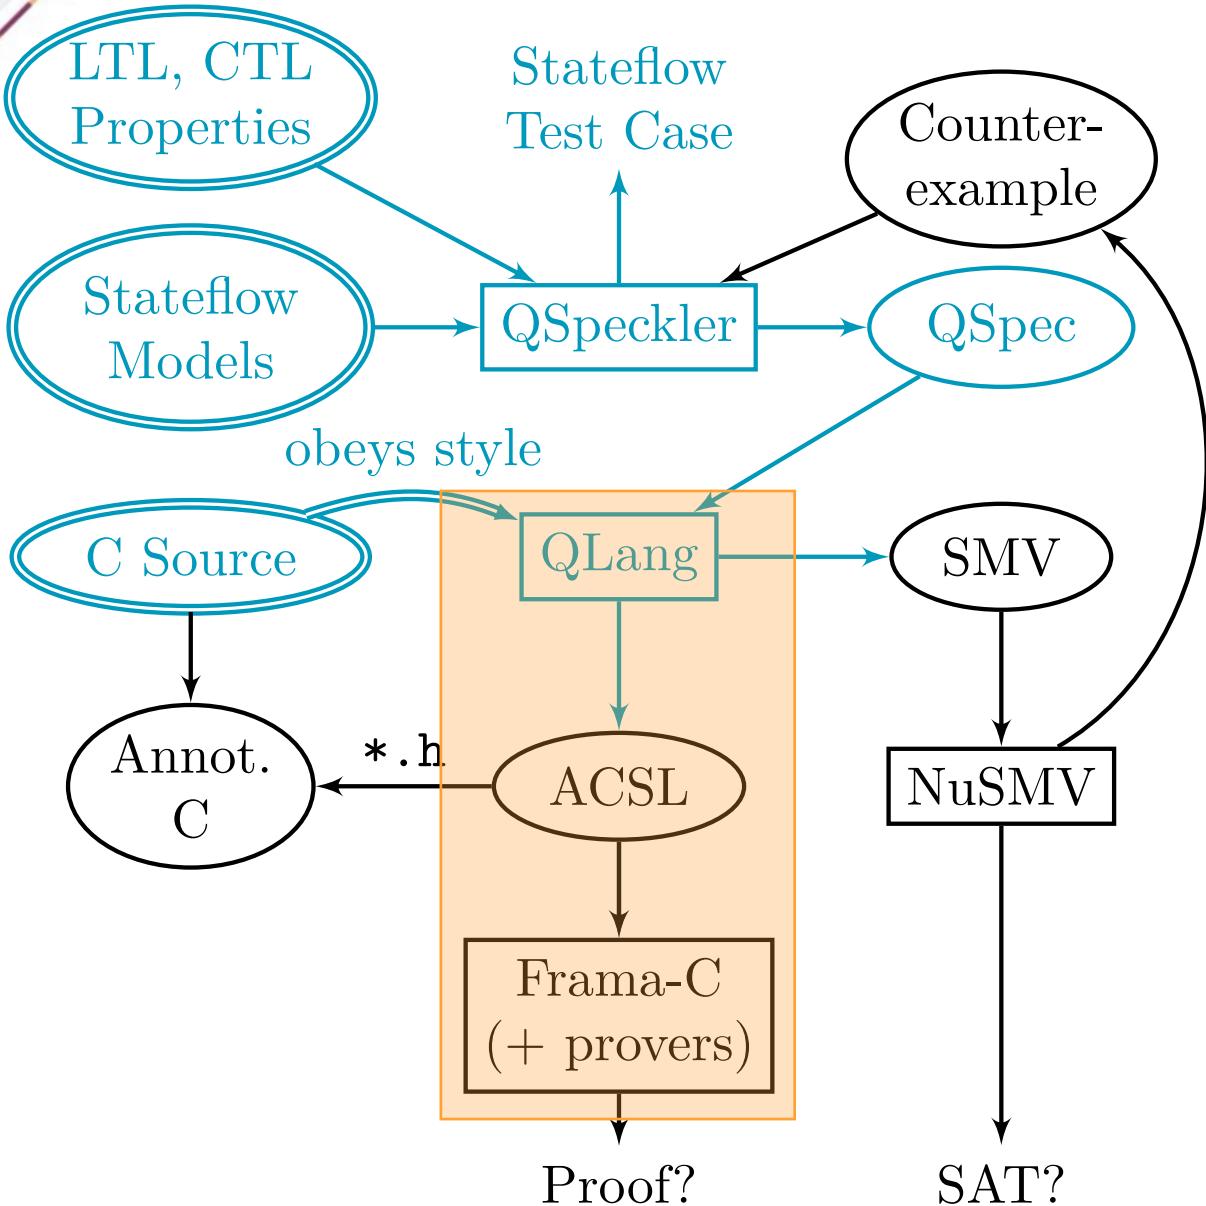
Modeling a Simple Clock



\leq weak



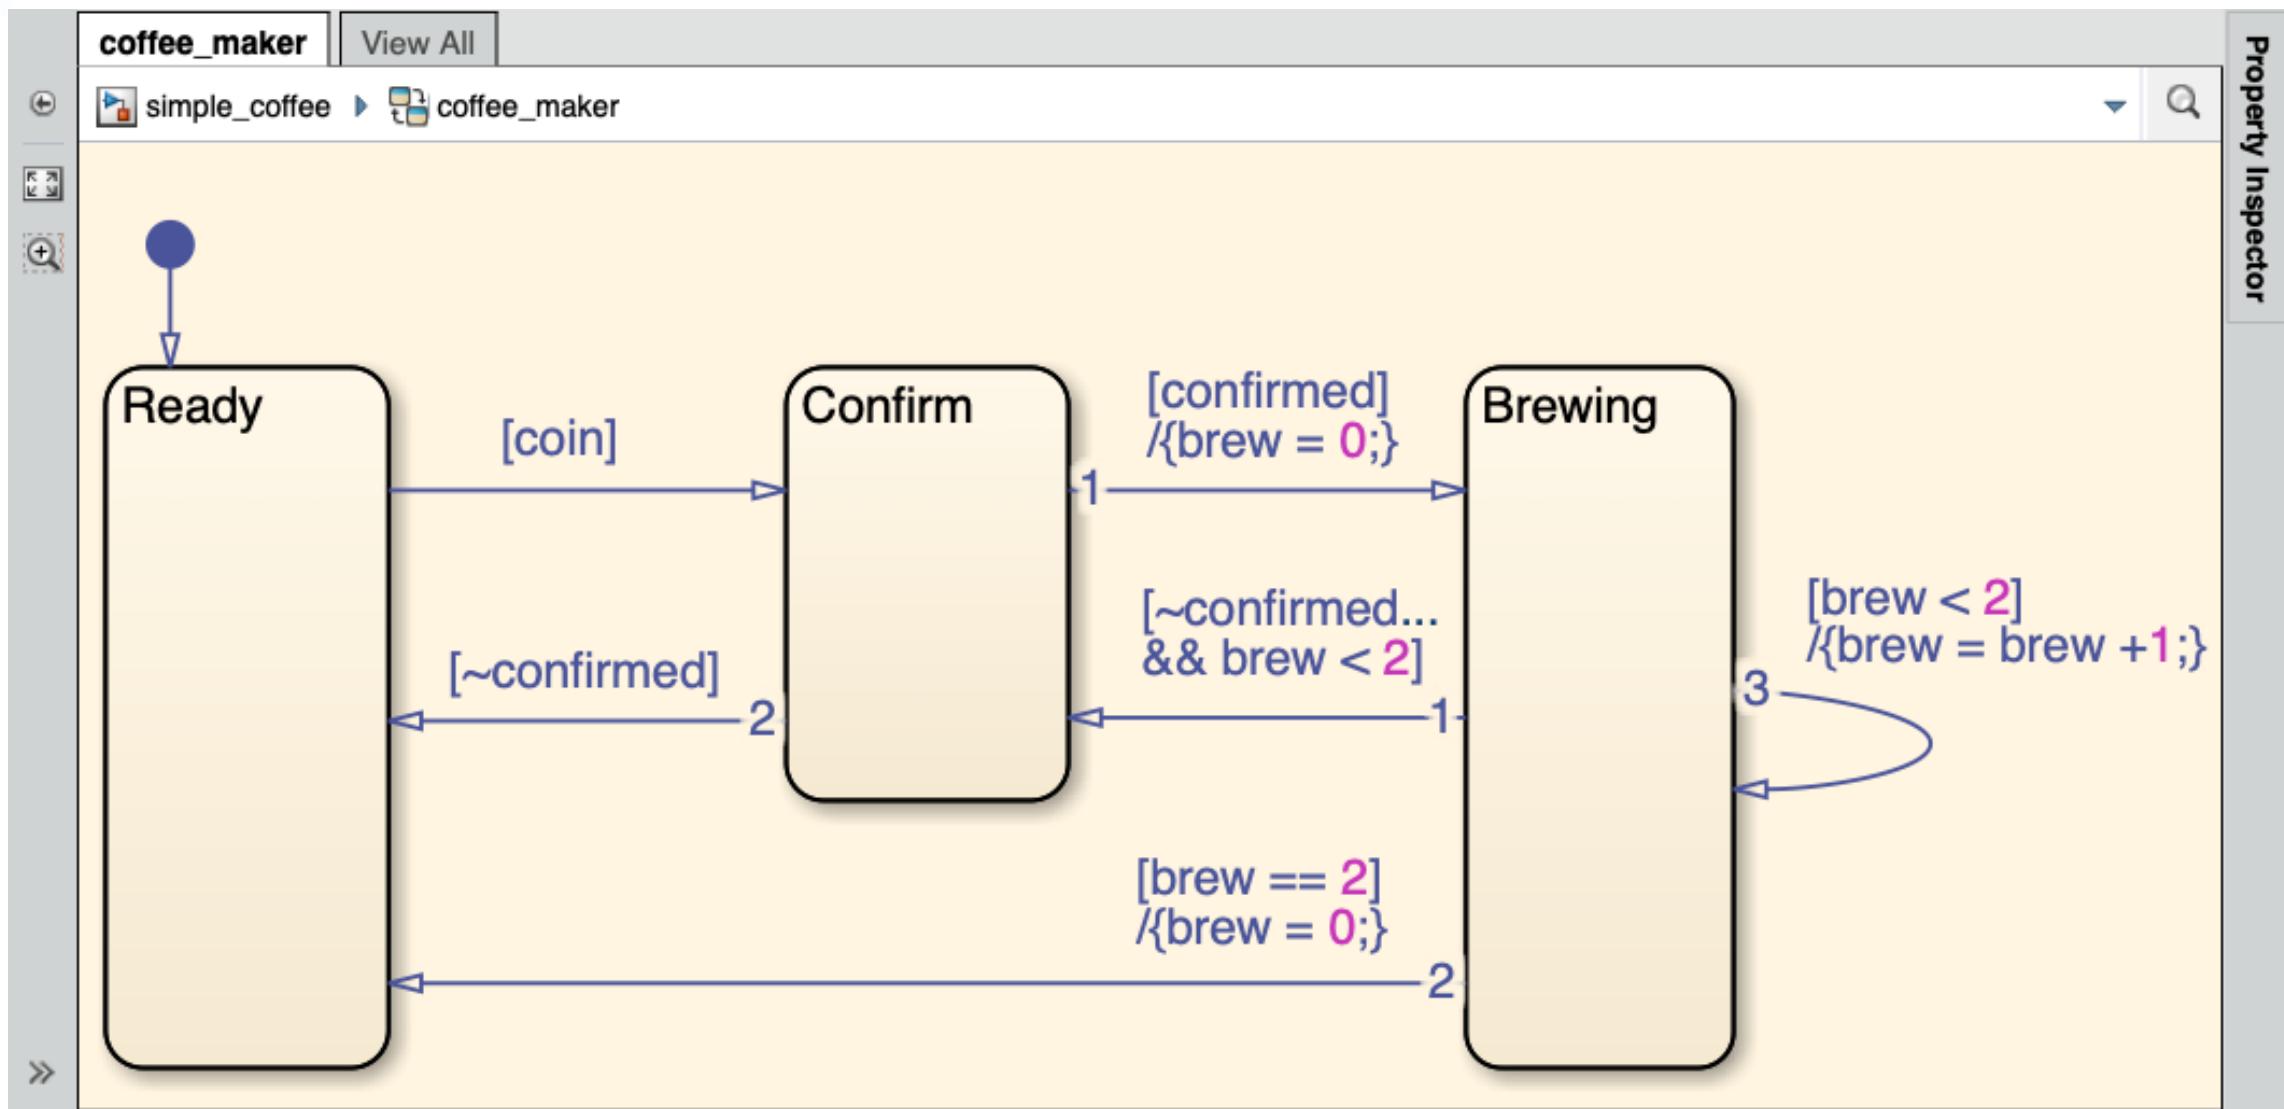
Architecture of Q Framework



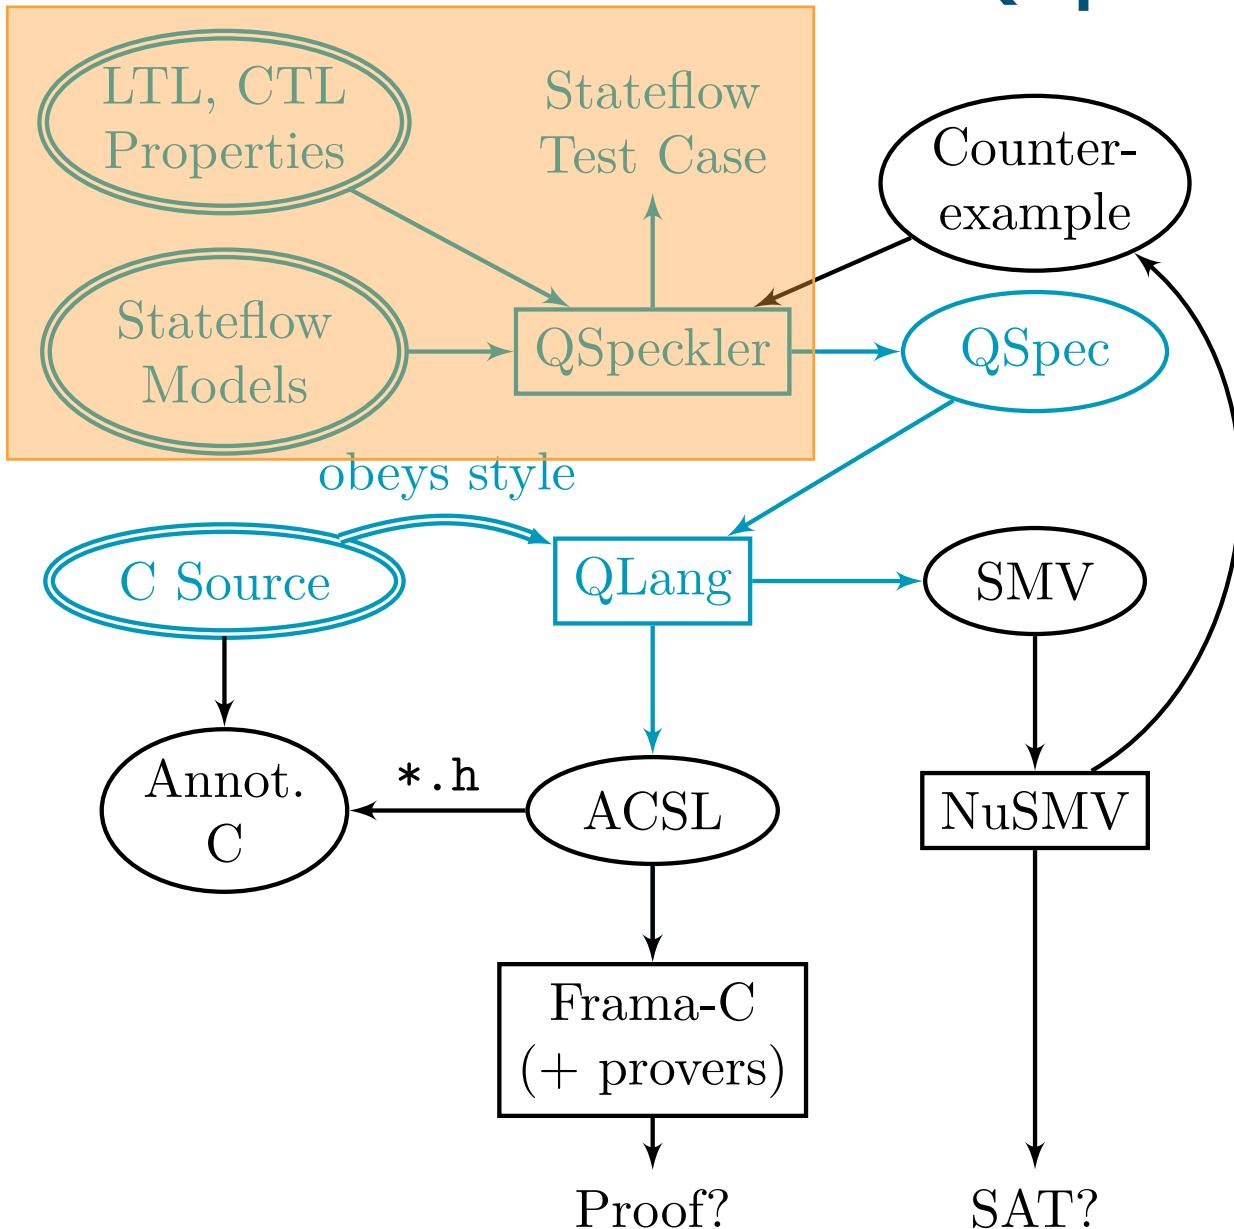
- Blue text: Sandia developed
- Double-struck: Written or checked by hand



Stateflow



Convert Stateflow to QSpec



- MATLAB App to generate SC-XML
- MATLAB expression parser
- Convenient UI for testing



QSpec

```
<?xml version="1.0" encoding="UTF-8"?>
<qspec>
    <!-- ... other initialization ... -->
    <sequential id="Clock">
        <variable id="tick" domain="boolean" intent="input"/>
        <variable id="h" domain="(range 0 23)" intent="register"/>
        <plain id="H"/>
        <transition type="initial" target="H">
            <assign id="h" ex="0"/>
        </transition>
        <transition source="H" target="H">
            <guard ex="tick"/>
            <assign id="h" ex="(ite (= h 23) 0 (+ h 1))"/>
        </transition>
    </sequential>
</qspec>
```

- Based on SCXML



Preliminaries

- A labeled transition system (LTS) is a triple
 (S, O, \rightarrow)
states, observations (labels), transition relation
- We are building a refinement between two LTSes
 $P_C \leqslant_{\text{weak}} Q$
 - P_C is a C program
 - Q is a QSpec
- Provided we can think of a C program as an LTS



Preliminaries

To define refinement, we first define partial correctness:

$$\{p\}f\{q\} := \forall s \in \text{ProgState}. \quad (1)$$

$$s \models p \implies (\forall s' \in \text{ProgState}. s \llbracket f \rrbracket s' \implies s' \models q),$$

WP's Hoare logic and predicate transformer semantics $\llbracket \cdot \rrbracket$

- But for Labeled Transition Systems, correctness is *stuttering-invariant trace equivalence*.



Comparing an LTS with C

- Strict refinement too strong
- Consider

$$\{p\}f\{q\}$$

- Frama-C cannot describe intermediate states
- Gives us modularity, but not observational refinement



Observable Events in C

- We require observational refinement
- We borrow CompCert's notion
 - externally-visible reads and writes
- Nontermination not included here
 - Design requirement
 - infinite event loop with handler
 - handlers are loop free

```
struct machine;
while(true) {
  msg = read_msg();
  if (msg == A)
    handle_A(&m);
  else
    handle_other(&m);
}
```

Weak Simulation

- So, we map observables into transitions in the LTS:

$$P \leq_{weak} Q := \forall(p, q) \in R, \alpha \in O_P, p' \in S_P. \quad (3)$$

$$p \xrightarrow{\alpha} p' \implies \exists q' \in S_Q. \left(q \xrightarrow{\tau^*} \xrightarrow{\alpha} \xrightarrow{\tau^*} q' \wedge (p', q') \in R \right),$$

- τ is the silent transition
- S the set of states
- $R \subseteq S_P \times S_Q$
- O an observable (*Label* in typical LTL notation)

Handling Volatile Reads and Writes

- Require any access wrapped in a function call
- Axiomatize hardware access
- Use *ghost state*

```
/*@  
ghost int obs_t;  
axiomatic model {  
    type obs;  
    logic obs obs_at(integer t);  
    logic uint8_t fgetC0bs(obs o);  
} */  
volatile uint8_t fgetCVal;
```

Weak Simulation



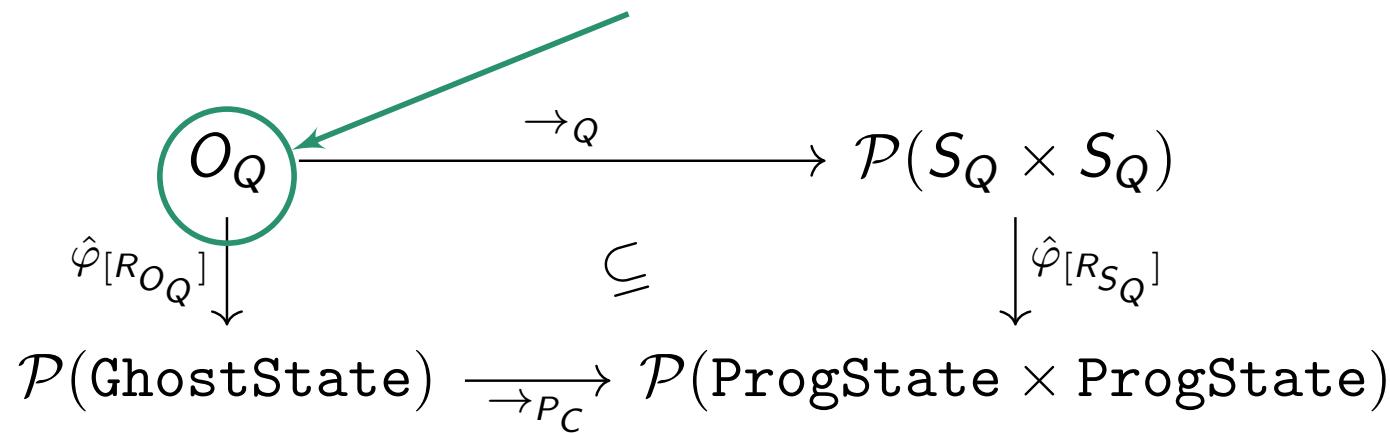
$$\begin{array}{ccc} O_Q & \xrightarrow{\rightarrow_Q} & \mathcal{P}(S_Q \times S_Q) \\ \hat{\varphi}_{[R_{O_Q}]} \downarrow & \subseteq & \downarrow \hat{\varphi}_{[R_{S_Q}]} \\ \mathcal{P}(\text{GhostState}) & \xrightarrow{\rightarrow_{P_C}} & \mathcal{P}(\text{ProgState} \times \text{ProgState}) \end{array}$$

- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

Weak Simulation

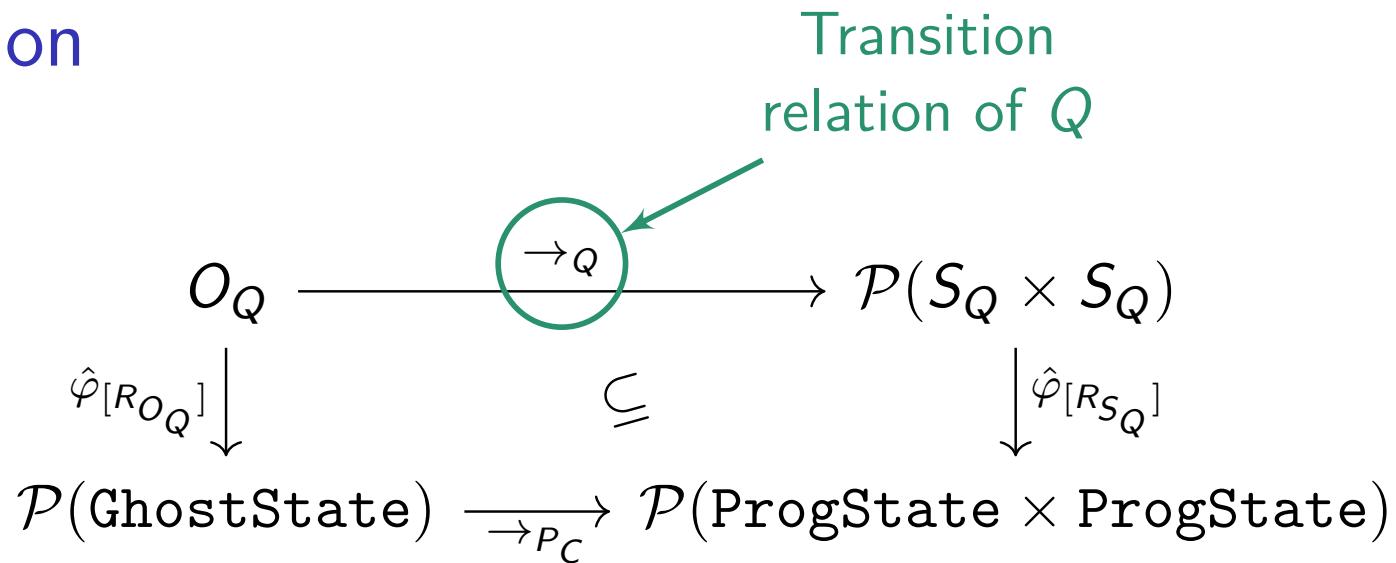


Observables in the LTS Q



- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

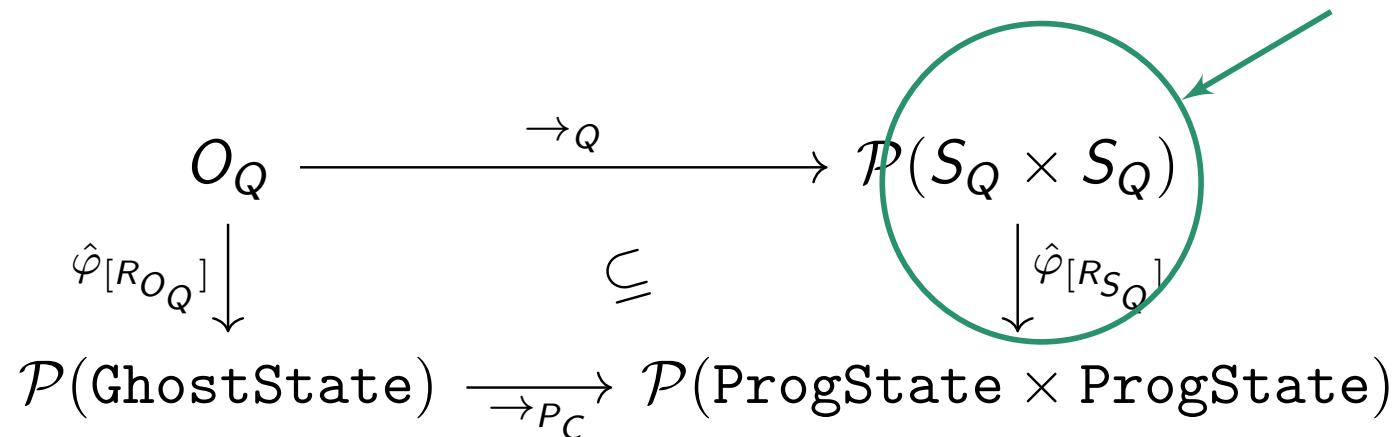
Weak Simulation



- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

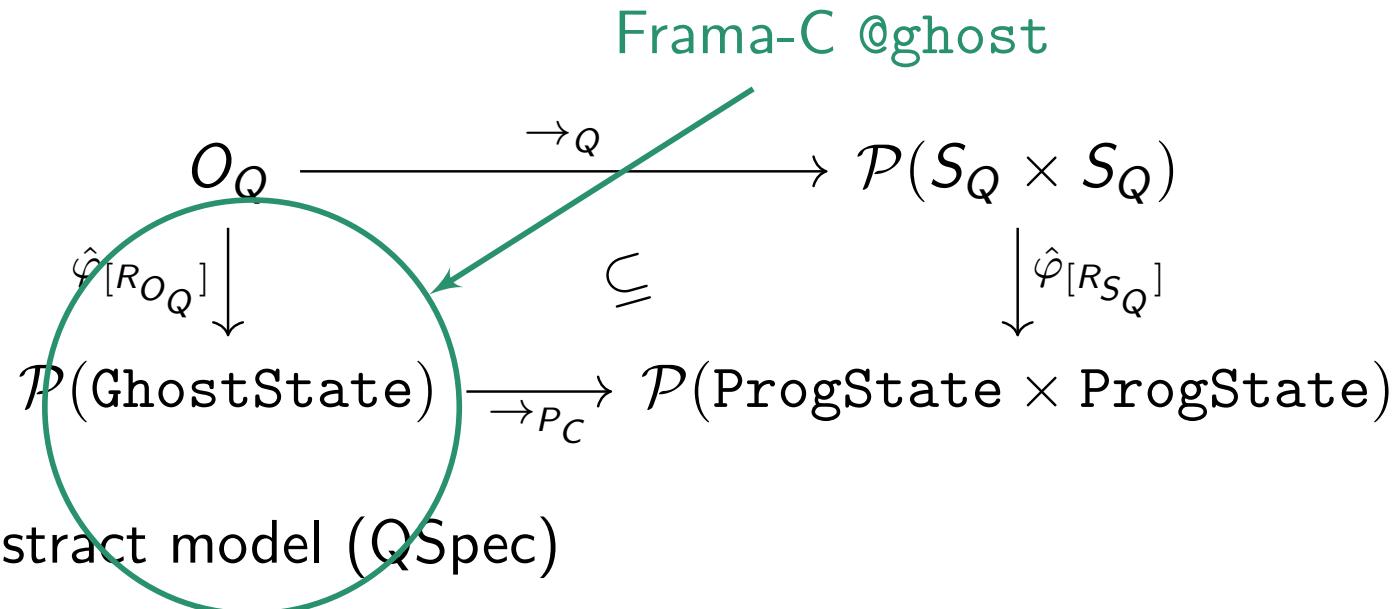
Weak Simulation

Relations over
states in Q



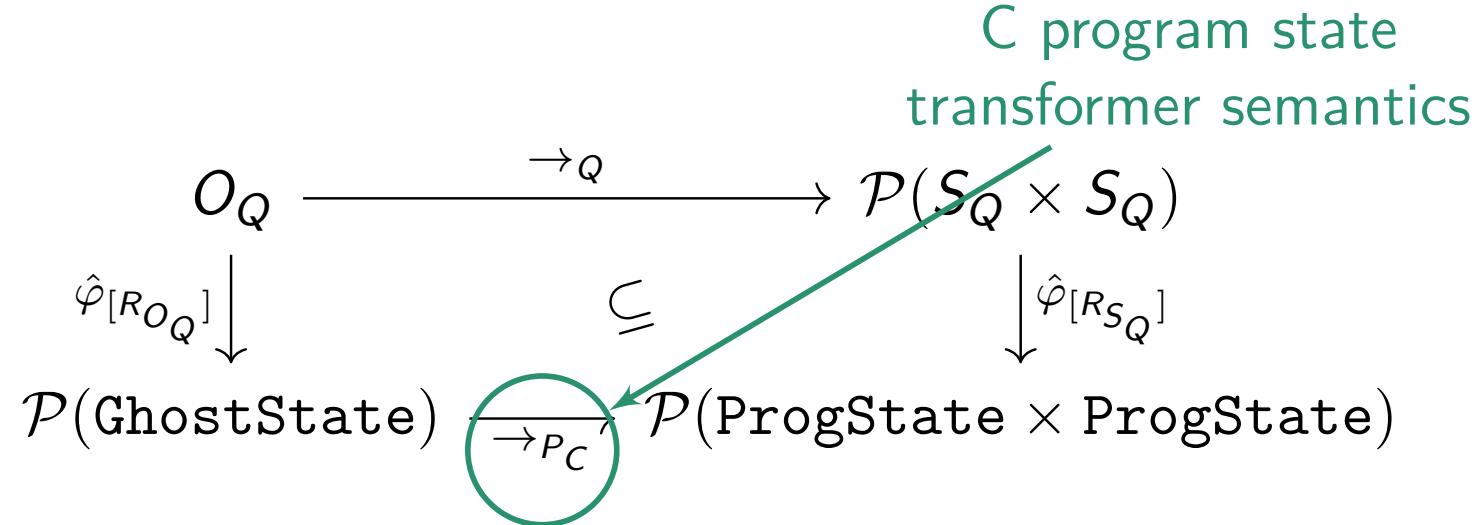
- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

Weak Simulation



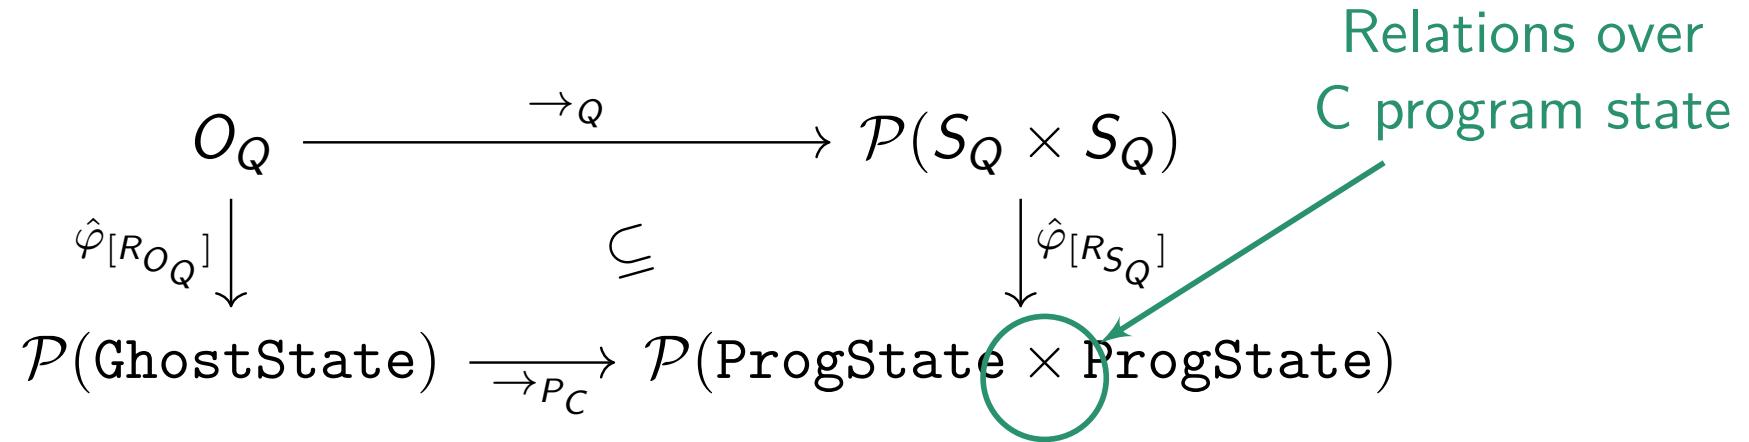
- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

Weak Simulation



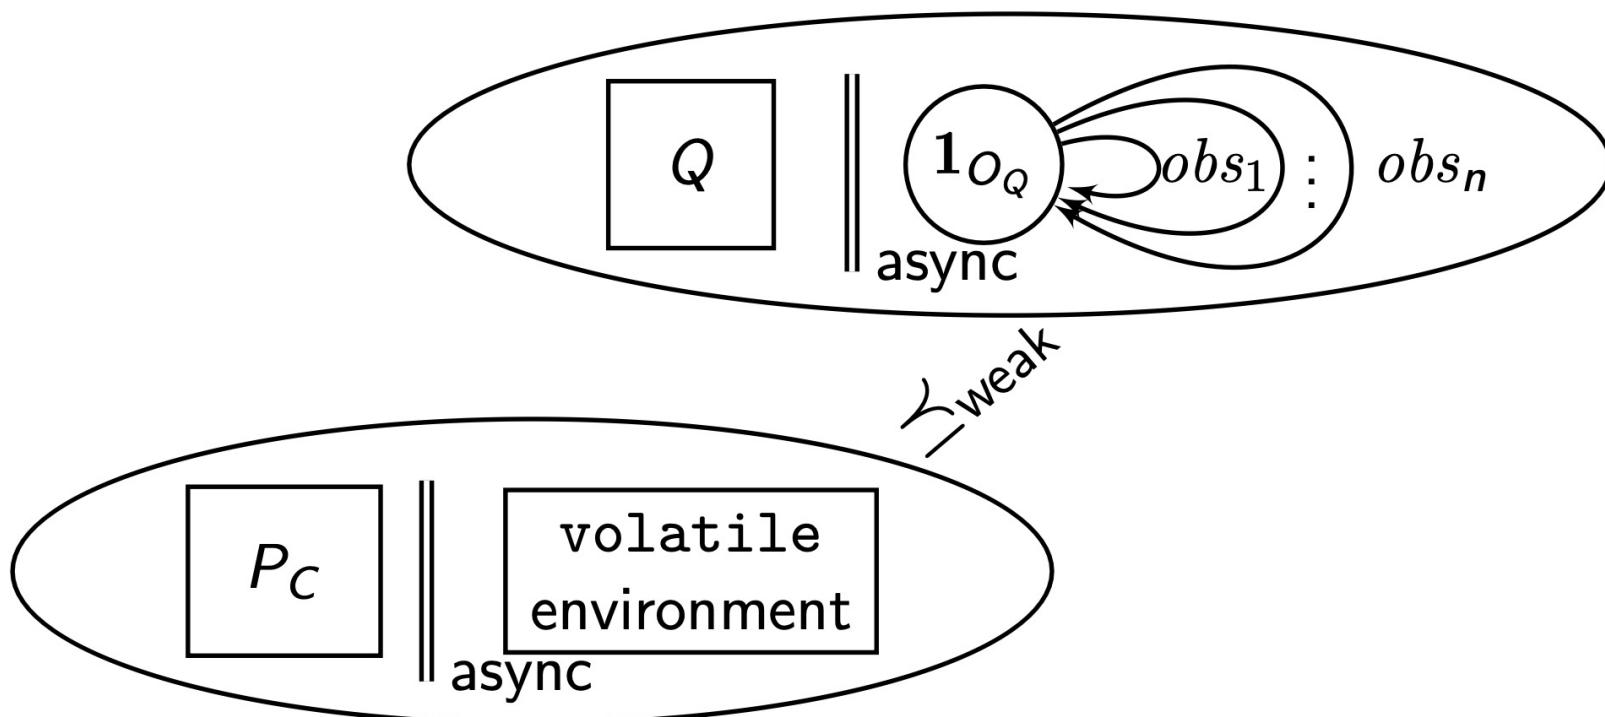
- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

Weak Simulation



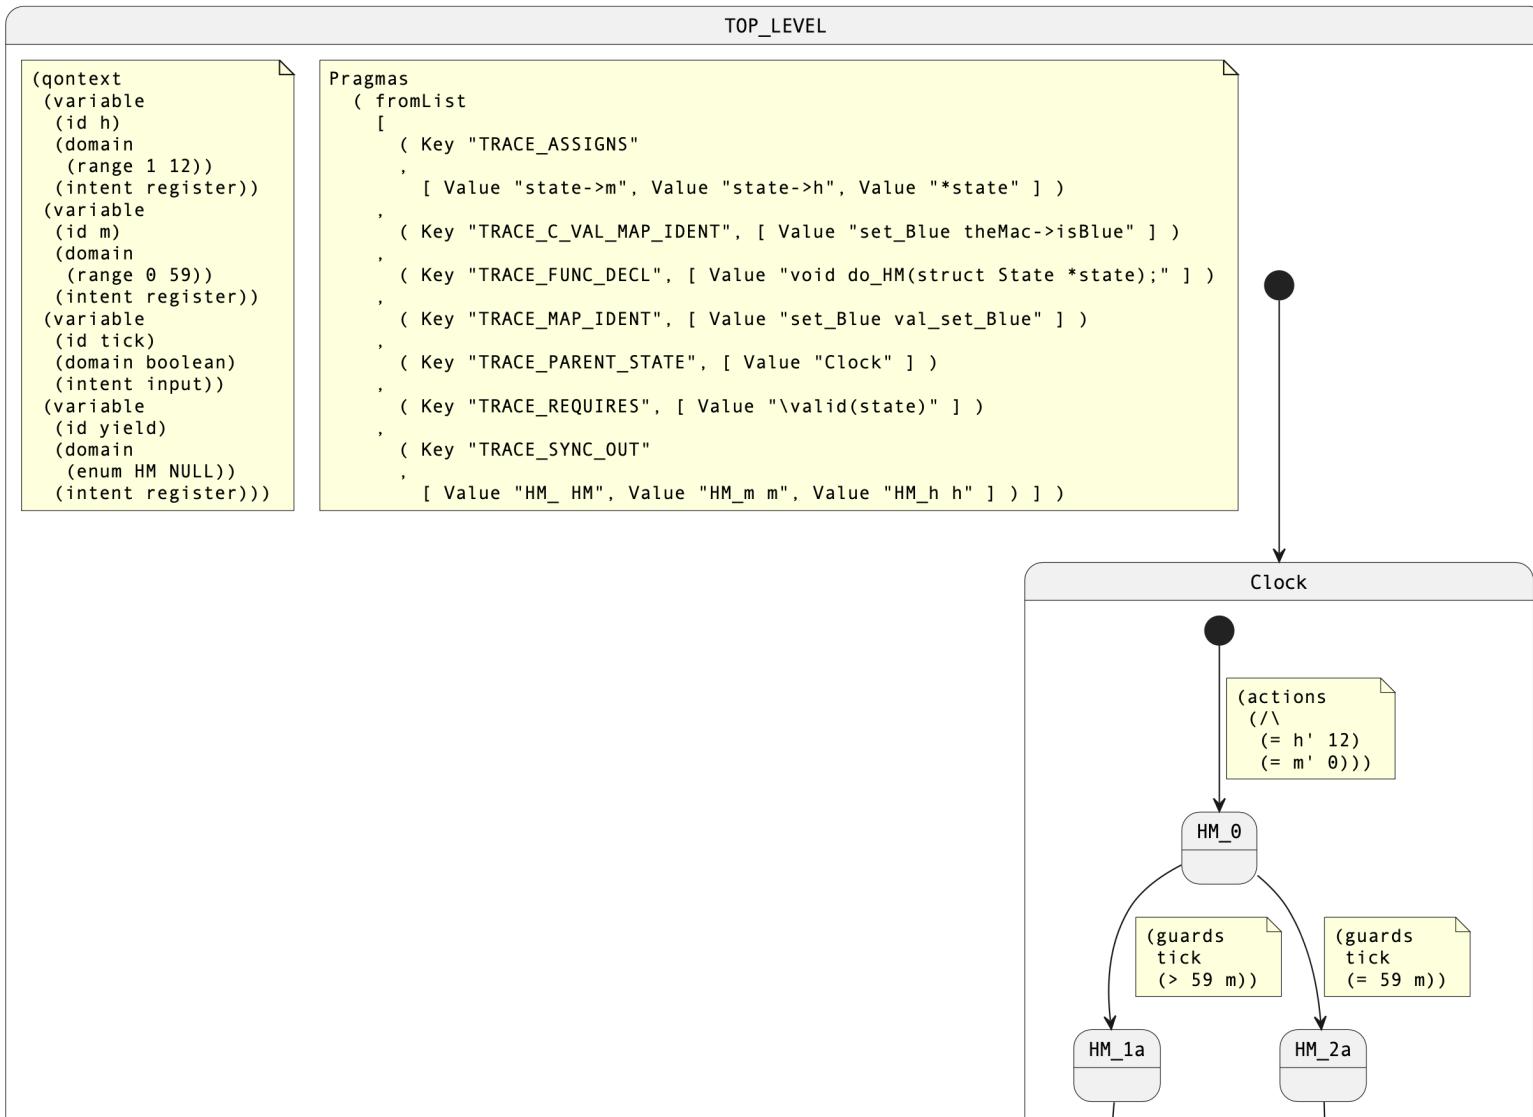
- Q is the abstract model (QSpec)
- P_C is the concrete implementation (C program)
- $\hat{\varphi}$ is a JSON file relating Stateflow variables to predicates over C variables.
- \rightarrow_Q is a Galois connection between O_Q and $\mathcal{P}(S_Q \times S_Q)$
- This demonstrates a proof of weak simulation, provided we can think of P_C as a transition system: this is not trivial when considering C semantics

Refinement



- Above: Composition of the model with an LTS with a single state 1
- Below: Composition in the C program with an environment for volatiles

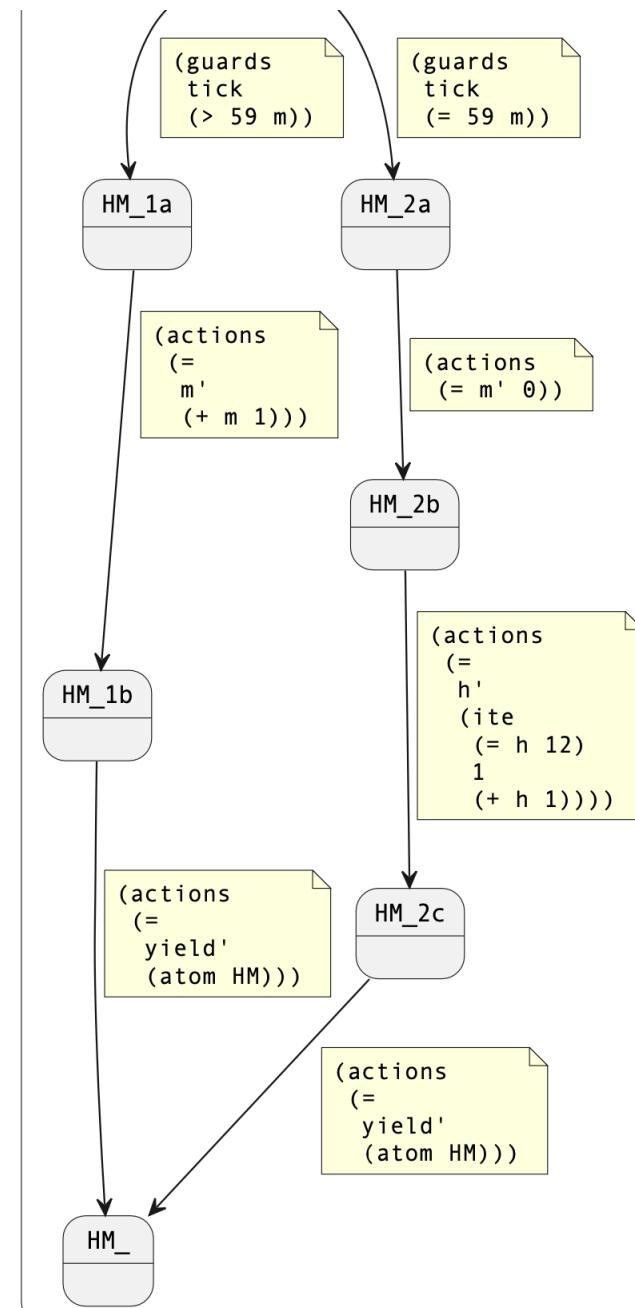
Example: Loop Free Machine



- Pragmas to link C with State Machine (Simulation Map)



Loop Free Machine



Generated ACSL

```
/* Generated with qlang */
/*@
axiomatic internal_states_Clock {
    logic gstate HM_1a;
    logic gstate HM_1b;
    logic gstate HM_2a;
    logic gstate HM_2b;
    logic gstate HM_2c;
}
predicate spec_step_Clock(integer t, integer ft) =
\let gs0 = gstate_at(t+0);
\let gs1 = gstate_at(t+1);
(gs0 == HM_0
  \&& m_at(t+1, 0).tick
  \&& (59 > m_at(t+1, 0).m)
  \&& ft == 0 ==> gs1 == HM_1a) \&&
    // ... other transitions */

/*@
// Behaviors -- all paths
behavior Path0000_Clock:
    assumes m_at(oracle_t+1, 0).tick
    \&& (59 > m_at(oracle_t+1, 0).m);
    ensures ...
behavior Path0001_Clock:
    ...
complete behaviors;
*/
void do_HM(struct State *state);
```



Voilà, the Trace Back-End

- Key idea
 - Enumerate all paths from initial state to terminals
 - Update ghost state, track all guards and actions along each path
 - Provided simulation map, this proves that C refines LTS
- Some Notes
 - Simulation Map can get complex; extra logic for:
 - handling nondeterminism (e.g., messages)
 - WP tactics
 - additional requires/ensures, error states,
 - Even so, most effort goes into generating & interpreting WP



Challenges

- Memory model
 - e.g., unions, bit-level operations
 - Granularity of assigns statements
- Counterexample generation
- Floating-point support is limited
- Scale: interpreting results from autogenerated proof obligations

```
Goal Check 'oracle_3' (file state.c, line 23):
Let a = L_m(oracle_t_0, ft_t_0).
Let a_1 = L_m(oracle_t_1, ft_t_3).
Let a_2 = shiftfield_F10_machine_nextState(theMac_0).
Let x = Mint_undef_0[a_2].
Let x_1 = Mint_0[a_2].
Let a_3 = shiftfield_F10_machine_currState(theMac_0).
Let m = Mint_0[a_3 <= 0].
Assume {
  Type: is_uint32_chunk(Mint_0) /\ is_bool(check_side_error_0) /\ 
        is_bool(old_val_bflushed_0) /\ is_bool(old_val_bit_delay_0) /\ 
        is_bool(old_val_dflushed_0) /\ is_bool(old_val_faultB_0) /\ 
        is_bool(old_val_faultD_0) /\ is_bool(old_val_side_err_0) /\ 
        is_uint32(old_val_set_Blue_0) /\ is_uint32(old_val_set_Green_0) /\ 
        is_uint32(old_val_set_Red_0) /\ is_sint32(ft_t_0) /\ 
        is_sint32(ft_t_1) /\ is_sint32(ft_t_2) /\ is_sint32(ft_t_3) /\ 
        is_sint32(ft_t_4) /\ is_sint32(ft_t_5) /\ is_sint32(old_t_plus_1_0) /\ 
        is_sint32(oracle_t_0) /\ is_sint32(oracle_t_1) /\ 
        is_sint32(oracle_t_2) /\ is_sint32(read_msg_if_ready_dev1_0) /\ 
        is_uint32_chunk(m) /\ is_uint32(x_1) /\ is_uint32(x) /\ 
        is_uint32_chunk(havoc(Mint_undef_0, m, theMac_0, 17)).
(* Heap *)
Type: (region(theMac_0.base) <= 0) /\ linked(Malloc_0).
(* Assertion 'rte,mem_access' *)
Have: valid_rw(Malloc_0, a_3, 1).
(* Call 'periodic_msg' *)
Have: (x = x_1) /\ valid_rw(Malloc_0, theMac_0, 17).
(* Call 'sync_code' *)
Have: P_sync_t(old_val_bit_delay_0, old_val_side_err_0,
              old_val_set_Green_0, old_val_set_Red_0, old_val_set_Blue_0,
              old_val_dflushed_0, old_val_bflushed_0, old_val_faultD_0,
              old_val_faultB_0, ft_t_4, oracle_t_1, oracle_t_2, ft_t_5, 15).
(* Call 'check_side_error' *)
Have: ((a_1.F11_call_side_err) != 0) /\ 
      (((a_1.F11_val_side_err) != 0) <-> (check_side_error_0 != 0)) /\ 
      P_sync_ft(old_val_bit_delay_0, old_val_side_err_0, old_val_set_Green_0,
                old_val_set_Red_0, old_val_set_Blue_0, old_val_dflushed_0,
                old_val_bflushed_0, old_val_faultD_0, old_val_faultB_0, ft_t_3,
                oracle_t_1, ft_t_4, 7).
(* Call 'read_msg_if_ready_dev1' *)
Have: ((L_m(1 + oracle_t_1, 0).F11_is_ready_dev1) != 0) /\ 
      P_sync_t2(old_val_bit_delay_0, old_val_side_err_0, old_val_set_Green_0,
                old_val_set_Red_0, old_val_set_Blue_0, old_val_dflushed_0,
                old_val_bflushed_0, old_val_faultD_0, old_val_faultB_0,
                old_t_plus_1_0, ft_t_2, oracle_t_0, oracle_t_1, ft_t_3, 4).
(* Then *)
Have: read_msg_if_ready_dev1_0 != 0.
(* Call 'msg_is_Red' *)
Have: ((L_m(oracle_t_0, ft_t_1).F11_msg) = 8) /\ 
      P_sync_t0(old_val_bit_delay_0, old_val_side_err_0, old_val_set_Green_0,
                old_val_set_Red_0, old_val_set_Blue_0, old_val_dflushed_0,
                old_val_bflushed_0, old_val_faultD_0, old_val_faultB_0, ft_t_1,
                oracle_t_0, ft_t_2, 3).
(* Call 'set_Red' *)
Have: ((a.F11_call_set_Red) != 0) /\ ((a.F11_val_set_Red) = 3) /\ 
      P_sync_ft(old_val_bit_delay_0, old_val_side_err_0, old_val_set_Green_0,
                old_val_set_Red_0, old_val_set_Blue_0, old_val_dflushed_0,
                old_val_bflushed_0, old_val_faultD_0, old_val_faultB_0, ft_t_0,
                oracle_t_0, ft_t_1, 0).
}
Prove: oracle_t_0 = 3.
Prover Z3 4.11.0 returns Timeout (Qed:26ms) (20s)
```



Future Work

- Open Source: currently in the process
- Formalization in Coq
 - Some parts are proven in Coq
 - Want a formal proof of refinement
 - Composition: have parallel async, want nested composition
- Extend Hoare logic to better handle LTS to C refinement
- Check <https://proof.sandia.gov/> for updates
- Thank you!

Bonus Content

Other Uses of Frama-C at Sandia

- Verification of a Kalman Filter
- Separate out:
 1. memory safety
 2. numerical properties
 3. concurrency/ scheduling
- (1) good for Frama-C
- less so (2) and (3)

